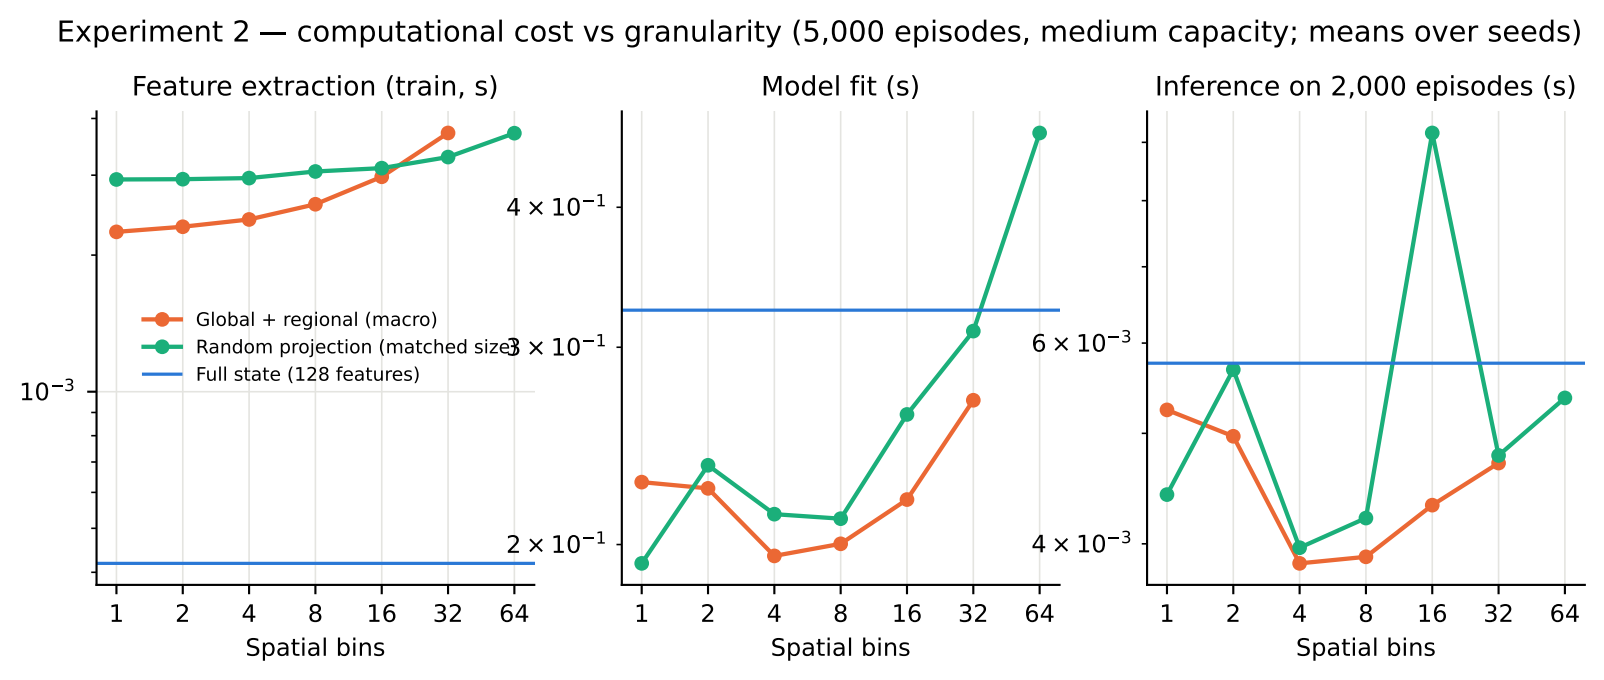
Experiment 2 — computational cost vs granularity (5,000 episodes, medium capacity; means over seeds)
Feature extraction (train, s)
Model fit (s)
Inference on 2,000 episodes (s)
Global + regional (macro)
Random projection (matched size)
Full state (128 features)
10^{-3}
2×10^{-1}
3×10^{-1}
4×10^{-1}
4×10^{-3}
6×10^{-3}
1
2
4
8
16
32
64
Spatial bins